

# Lab 4: Transformers and Dyck Languages

ZHU Yixuan

## 4. Theoretical Aspects

### 1. The Chomsky hierarchy and Dyck languages

The Chomsky hierarchy classifies formal languages according to the computational power needed to recognise or generate them:

| Class  | Grammar type                     | Recogniser              |
|--------|----------------------------------|-------------------------|
| Type 3 | Regular languages                | Finite-state automata   |
| Type 2 | Context-free languages           | Pushdown automata       |
| Type 1 | Context-sensitive languages      | Linear-bounded automata |
| Type 0 | Recursively enumerable languages | Turing machines         |

Dyck languages are context-free languages. For example, the one-bracket Dyck language  $D(1)$ , containing strings such as  $()$ ,  $(( ))$ , and  $()()$ , is generated by the context-free grammar

$S \rightarrow \epsilon \mid SS \mid (S)$

For  $k$  bracket types, the grammar becomes

$S \rightarrow \epsilon \mid SS \mid a_i S a_{i\_bar} \quad \text{for } 1 \leq i \leq k$

This grammar is context-free because every production rewrites a single nonterminal  $S$  independently of its surrounding context. Dyck languages are generally not regular, because a finite-state automaton cannot keep track of an unbounded number of unmatched opening brackets. Recognising a Dyck string requires stack-like memory: every opener must be pushed, and every closer must match the most recent unmatched opener. This is exactly the kind of memory provided by a pushdown automaton.

$D(2)$  has a special theoretical role: every context-free language can be obtained as a homomorphic image of the intersection between  $D(2)$  and some regular language  $R$ . Informally, this means that the two-type Dyck language is expressive enough to encode the derivational structure of any context-free grammar, while the regular language  $R$  filters out only the well-formed encodings that correspond to valid derivations of the target grammar. A homomorphism then maps those encoded bracket structures back to the terminal strings of the original language.

So the construction has three parts:

1.  $D(2)$  supplies the universal nested, stack-like structure.
2. A regular constraint  $R$  selects the encodings that correspond to the desired context-free grammar.
3. A homomorphism erases or renames symbols to recover the original context-free language.

This makes Dyck languages a natural benchmark for structural generalisation. They isolate one of the central difficulties of context-free syntax, that is unbounded hierarchical dependency. A model that succeeds only on short or shallow Dyck strings may have learned surface statistics, such as frequent local bracket patterns. A model that succeeds on deeper nesting than it saw during training has stronger evidence of learning a more rule-like stack discipline. In this lab, this is exactly why training on  $D(2, n \leq 4)$  and testing on  $D(2, n \in \{5, 6, 7\})$  is informative: the OOD split tests whether the Transformer has learned a reusable structural rule or merely memorised the training distribution.

## 2. Concrete examples

**Example 1: Parentheses, brackets, and braces in source code** Programming languages contain many nested delimiter structures, for example, parentheses in function calls, square brackets in indexing, braces in blocks, and angle brackets in some type systems. If we abstract away from the internal tokens and keep only the delimiters, the well-formedness condition is essentially a Dyck language with several bracket types.

For example, the Python expression

```
result = f(arr[i], (x + y))
```

has the delimiter projection

```
( [ ] ( ) )
```

which is a well-balanced Dyck string. Similarly, in C-like languages, nested blocks and function calls produce strings over bracket types such as `()`, `[]`, and `{}`.

The four error types correspond naturally to common programming mistakes:

| Dyck error type    | Software-engineering analogue                                                             |
|--------------------|-------------------------------------------------------------------------------------------|
| E1 Missing closer  | Forgetting a closing parenthesis, bracket, or brace, e.g. <code>f(x, y</code>             |
| E2 Spurious opener | Adding an extra opening delimiter, e.g. <code>f((x, y</code>                              |
| E3 Type mismatch   | Closing with the wrong delimiter, e.g. <code>arr[i)</code> instead of <code>arr[i]</code> |
| E4 Premature close | Writing a closing delimiter before anything has opened, e.g. <code>)x + y(</code>         |

In real compilers and IDEs, these errors are often diagnosed by a stack-like parser. Openers are pushed onto a stack, and closers must match the top of the stack.

**Example 2: Nested markup in HTML or XML** Markup languages such as HTML and XML also have nested paired delimiters. If we treat an opening tag as an opening bracket and its corresponding closing tag as the matching closing bracket, then well-formed documents again reduce to a Dyck-like language. For example,

```
<div><em>text</em><span>more</span></div>
```

has the abstract structure

```
div_open em_open em_close span_open span_close div_close
```

which behaves like a multi-type Dyck string. The difference from the toy setting is that the number of possible tag types can be large, but the basic stack discipline is the same.

The error types have direct markup analogues:

| Dyck error type    | Markup analogue                                                                                            |
|--------------------|------------------------------------------------------------------------------------------------------------|
| E1 Missing closer  | Missing a closing tag, e.g. <code>&lt;div&gt;&lt;em&gt;text&lt;/em&gt;</code>                              |
| E2 Spurious opener | Extra unclosed opening tag, e.g. <code>&lt;div&gt;&lt;span&gt;&lt;em&gt;text&lt;/em&gt;&lt;/div&gt;</code> |
| E3 Type mismatch   | Closing the wrong tag, e.g. <code>&lt;div&gt;&lt;em&gt;text&lt;/div&gt;&lt;/em&gt;</code>                  |
| E4 Premature close | A closing tag appears before its opener, e.g. <code>&lt;/span&gt;&lt;div&gt;text&lt;/div&gt;</code>        |

This example is especially close to the Dyck task because validating nested tags requires remembering not only how many constituents are open, but also their types and order. A model must therefore learn more than a simple count, that is to say, it must learn a last-in-first-out matching relation.

**Alternative computational-linguistics example: constituency syntax** Constituency parse trees can also be represented with brackets. For instance,

(S (NP the dog) (VP chased (NP the cat)))

is well-formed only if every opened constituent is eventually closed in the correct order. A missing closing bracket corresponds to an incomplete constituent; a spurious opener corresponds to an extra phrase boundary; a type mismatch corresponds to ending the wrong syntactic category; and a premature close corresponds to closing a constituent that was never opened. Natural language syntax is not just Dyck structure, but Dyck languages capture the nested dependency component that parsers must handle.

### 3. (bonus) LSTMs, memory, Transformers for Dyck languages

For  $D(1)$ , there is only one bracket type. Recognising a string only requires checking a counter:

1. Increment the counter when an opener ( is read.
2. Decrement it when a closer ) is read.
3. Reject if the counter ever becomes negative.
4. Accept if the final counter is zero.

If the maximum nesting depth is bounded by  $n$ , then the counter only needs to represent values from 0 to  $n$ . In practice, an LSTM with hidden dimension  $h$  can encode this bounded counter in its continuous hidden state, at least for depths within the range seen during training. This could explain why  $D(1)$  is relatively easy. the model does not need to remember a sequence of bracket types, only the current depth.

For  $D(k)$  with  $k \geq 2$ , counting is no longer sufficient. The model must remember the type of each currently unmatched opener, in last-in-first-out order. For example,  $([])$  and  $[()]$  may have the same depth profile, but they require different future closing brackets. At depth  $n$ , there are  $k^n$  possible stacks of opener types. A recogniser must distinguish among these possible stacks, because the correct next closer depends on the top of the stack, and later closers depend on the lower stack elements.

Formally, consider prefixes consisting of  $n$  openers over  $k$  types. There are  $k^n$  different prefixes:

$a_{i1} a_{i2} \dots a_{in}$

Each prefix determines a unique sequence of valid closers:

$\text{abar}_{in} \dots \text{abar}_{i2} \text{abar}_{i1}$

If two different opener stacks were mapped to the same internal memory state, the recogniser would have to give the same decision for all possible continuations from that state. But because the two stacks differ at some position, there exists a closing continuation that is valid for one stack and invalid for the other. Therefore, an exact recogniser needs enough memory states to distinguish all possible stacks up to depth  $n$ . The amount of information needed grows as  $\log_2(k^n) = n \log_2(k)$  bits. Thus, for  $k \geq 2$ , the memory requirement grows linearly with the nesting depth  $n$ .

An LSTM with fixed hidden dimension has fixed-size continuous memory. In theory, real-valued states can encode very large amounts of information with arbitrary precision, but in realistic finite-precision computation and finite-data training, this is not a reliable unbounded stack. So an LSTM may learn to recognise  $D(k)$  for bounded depths, but its ability to generalise to larger unseen depths is limited by whether it has learned a true stack-like algorithm rather than a depth-specific encoding.

For a Transformer with fixed width and fixed number of layers, the situation is also constrained. Self-attention gives every token direct access to every other token, which is useful for bracket matching. However, the model still has fixed-dimensional token representations and a fixed number of computation steps. If training only covers depths up to  $n \leq 4$ , the model may learn attention patterns that work well for shallow examples, such as attending to nearby matching brackets or memorising frequent local configurations. Nothing in a standard fixed-width Transformer forces it to implement an unbounded pushdown stack.

This implies that OOD generalisation to deeper nesting is not guaranteed. If the Transformer has learned a reusable rule, performance should degrade little when moving from  $n \leq 4$  to  $n = 5, 6, 7$ . If it has learned shortcuts tied to the training distribution, performance should drop as depth and length increase. The OOD experiments in this lab are therefore a test of whether the model's internal representations behave like an algorithmic stack or like a collection of bounded-depth pattern recognisers.

## 5. Training and in-distribution evaluation

### 4. Train a binary classifier with [CLS] and error detection.

I implemented the binary detection model in `lab4_detection.py`. The script contains:

- synthetic Dyck string generation for the two bracket types `()` and `[]`;
- the four corruption operations E1-E4;
- tokenisation with `[CLS]`, `[SEP]`, and `[PAD]`, padded to length 80;
- a small BERT-style Transformer encoder;
- binary classification from the final `[CLS]` representation;
- accuracy, macro-F1, training curves, and confusion matrices.

For detection, I used a balanced binary split, 50% correct strings and 50% corrupted strings. The corrupted half is itself balanced across E1-E4. This avoids the trivial majority-class solution that would arise if the five descriptive categories were all equally frequent, since that would make 80% of the examples erroneous.

The model architecture was:

| Hyperparameter          | Value      |
|-------------------------|------------|
| Layers                  | 2          |
| Attention heads         | 4          |
| Hidden dimension        | 64         |
| Feed-forward dimension  | 256        |
| Dropout                 | 0.1        |
| Positional encoding     | Sinusoidal |
| Maximum sequence length | 80         |

The lab handout recommends 50,000 training examples and 5,000 examples for development and test. Because of time and compute constraints, the experiments in this report use a reduced setup: 10,000 training examples, 1,000 development examples, and 1,000 test examples for the main detection run. The results should therefore be read as preliminary but informative, rather than as a full-scale reproduction of the recommended setting.

The detection experiment was run as follows:

```
python3 lab4_detection.py \  
  --train-size 10000 \  
  --dev-size 1000 \  
  --test-size 1000 \  
  --epochs 8 \  
  --batch-size 128 \  
  --output-dir artifacts/detection
```

The test results were:

| Metric    | Value |
|-----------|-------|
| Accuracy  | 0.852 |
| Macro-F1  | 0.849 |
| Test loss | 0.358 |

The binary confusion matrix was:

|               | Predicted no error | Predicted error |
|---------------|--------------------|-----------------|
| Gold no error | 497                | 3               |
| Gold error    | 145                | 355             |

Broken down by error type:

| Error type         | Predicted no error | Predicted error |
|--------------------|--------------------|-----------------|
| No error           | 497                | 3               |
| E1 Missing closer  | 11                 | 114             |
| E2 Spurious opener | 12                 | 113             |
| E3 Type mismatch   | 14                 | 111             |
| E4 Premature close | 108                | 17              |

The easiest errors were E1, E2, and E3. E1 and E2 change the length parity of the sequence: a well-formed Dyck string must have even length, while deleting one closer or inserting one opener creates an odd-length sequence. E3 introduces a relatively local inconsistency, since a closer of the wrong type conflicts with the expected opener type.

E4 was systematically harder in this run, but this should not be explained by the absence of a length-parity cue. E4 is also implemented as an insertion, so it changes the length parity just as E2 does. The low E4 detection rate therefore suggests either that the model did not use the parity cue reliably, or that this particular E4 generation/distribution deserves additional inspection. As a data sanity check, the generated E4 examples were also evaluated with the deterministic PDA baseline and were correctly classified as invalid in that evaluation. Still, this result should be interpreted cautiously rather than as evidence that E4 is intrinsically harder because it lacks a simple length cue.

I initially expected E4 to be easy, since a premature closer violates prefix well-formedness. The confusion matrix showed the opposite, which made me check whether the model was really exploiting simple length or parity cues. The answer seems to be: not reliably.

The training curves and confusion matrix were saved to:

```
artifacts/detection/training_curves.png
artifacts/detection/confusion_by_error_type.png
artifacts/detection/metrics.json
```

The model improved substantially late in training: development accuracy moved from approximately 0.51 after the first epoch to 0.86 after the eighth epoch. There is no strong evidence of over-fitting in this preliminary run, since development loss continued to decrease by the final epoch. The regularisation used here was dropout 0.1, AdamW weight decay 0.01, and gradient clipping with maximum norm 1.0.

## 5. Error correction as sequence labelling

I implemented the local correction formulation in `lab4_correction.py`. Each input is the possibly corrupted sequence, and each token receives one edit label. I used the following label set:

```
OK
DELETE
INSERT_x
REPLACE_x
```

where `x` is one of the bracket tokens. In the current corruption scheme, only a subset of these labels is actually needed: missing closers require `INSERT_)` or `INSERT_]`; type mismatches require `REPLACE_)` or `REPLACE_]`; spurious openers and premature closers require `DELETE`.

The precise local convention is:

- `OK`: keep the current token;
- `DELETE`: remove the current token;
- `REPLACE_x`: replace the current token with `x`;
- `INSERT_x`: insert `x` after the current token.

For E1, where a closer has been deleted, the insertion label is placed on the token immediately before the missing closer. For E2 and E4, the inserted corrupted token is labelled `DELETE`. For E3, the mismatched closer is labelled `REPLACE_x`, where `x` is the original correct closer.

I trained the same small Transformer encoder as in the detection experiment, but with a token classifier on top of every contextual representation. Because `OK` labels are overwhelmingly more frequent than edit labels, I used inverse-square-root class weighting in the loss.




Figure 1: Detection confusion by error type

The experiment was run with:

```
python3 lab4_correction.py \
  --train-size 10000 \
  --dev-size 1000 \
  --test-size 1000 \
  --epochs 20 \
  --batch-size 128 \
  --weight-power 0.5 \
  --output-dir artifacts/correction
```

The test results were:

| Metric               | Value |
|----------------------|-------|
| Token-level accuracy | 0.909 |
| Exact-match accuracy | 0.509 |
| Test loss            | 0.311 |

Exact-match by type:

| Error type         | Exact-match accuracy |
|--------------------|----------------------|
| No error           | 1.000                |
| E1 Missing closer  | 0.000                |
| E2 Spurious opener | 0.008                |
| E3 Type mismatch   | 0.048                |
| E4 Premature close | 0.016                |

These results show why token-level accuracy is misleading for this task. Most tokens require no edit, so a model can obtain high token accuracy by predicting OK almost everywhere. Exact-match accuracy is much stricter: one wrong edit decision is enough to make the whole corrected string incorrect. The model preserves correct strings very well, but it does not yet reliably repair corrupted ones.

It is important to separate the no-error cases from the corrupted cases. The overall exact-match accuracy of 0.509 is mostly explained by the no-error examples, for which exact-match is 1.000. Since the dataset is balanced with 50% correct strings, an always-OK or no-edit baseline would already obtain about 0.50 exact-match accuracy by leaving all correct strings unchanged. On corrupted strings only, the model succeeds on only about 0.018 of examples:

| Subset                   | Exact-match accuracy |
|--------------------------|----------------------|
| All examples             | 0.509                |
| No-error examples        | 1.000                |
| Corrupted examples only  | 0.018                |
| Approx. no-edit baseline | 0.500                |

Thus, the correction experiment should be interpreted as a formulation and diagnostic baseline, not as a successful repair system. It shows that the local labelling setup can be implemented and evaluated, but the trained model almost never repairs corrupted strings exactly.

To show what a more appropriate correction procedure can achieve in this controlled setting, I also added a constrained single-edit repair baseline to `lab4_correction.py`. It enumerates all strings reachable by one edit from the corrupted input and keeps a candidate only if it is accepted by the deterministic Dyck PDA. This is not a neural sequence labeller, but it matches the structure of the lab task much better: there is exactly one corruption, and the corrected output should be a valid Dyck string.

| Method                 | Exact-match | Corrupted-only exact-match | Valid Dyck output rate |
|------------------------|-------------|----------------------------|------------------------|
| No-edit baseline       | 0.500       | 0.000                      | 0.500                  |
| Neural local labeller  | 0.509       | 0.018                      | not constrained        |
| PDA single-edit repair | 0.744       | 0.488                      | 1.000                  |

The PDA repair baseline is much closer to the ideal correction behaviour: every output is a valid Dyck string. Its exact match against the original clean string is lower than its validity rate because exact reconstruction is sometimes ambiguous. This is especially visible for E1 missing closers, where several insertion positions may restore a valid Dyck string even if only one matches the original pre-corruption string.

PDA single-edit exact-match by type:

| Error type         | Exact-match accuracy |
|--------------------|----------------------|
| No error           | 1.000                |
| E1 Missing closer  | 0.000                |
| E2 Spurious opener | 0.680                |
| E3 Type mismatch   | 0.464                |
| E4 Premature close | 0.808                |

PDA single-edit valid-output rate by type:

| Error type         | Valid Dyck output rate |
|--------------------|------------------------|
| No error           | 1.000                  |
| E1 Missing closer  | 1.000                  |
| E2 Spurious opener | 1.000                  |
| E3 Type mismatch   | 1.000                  |
| E4 Premature close | 1.000                  |

This makes the E1 result less surprising. For E1, the baseline always restores a valid Dyck string, but it does not know the original deletion position. Exact reconstruction can therefore fail even when the validity objective is satisfied.

This comparison changes how I interpret the correction part of the lab. The local sequence labelling model completes the requested formulation and evaluation, but it is not a good repair system in this small run. The constrained PDA repair baseline is the stronger correction result and should be used as the reference point for what is achievable when the grammar constraint is used explicitly.

This local sequence-labelling formulation has several limitations.

First, it is extremely class-imbalanced. In a sequence of length 40 with exactly one corruption, only one token usually has a non-OK label. As a result, token-level accuracy mostly measures whether the model learned to leave tokens untouched, not whether it learned to correct errors.

Second, local labels are not always a natural representation of a structural repair. A missing closer is an absence, not an observed token. Placing `INSERT_x` on the previous token is an implementation convention, but the true evidence for the missing closer may be distributed across the whole stack structure.

Third, the formulation predicts edits independently at each position. It does not enforce that the final edited string is a valid Dyck string, nor that there is exactly one correction. The model can therefore produce multiple incompatible edits, or no edit at all. A constrained decoder or a parser-based repair system would be better suited to enforcing global well-formedness.

Finally, exact correction may require distinguishing several globally plausible repairs. For example, inserting a missing closer at one position versus another can both improve local balance but produce different full strings. The local classifier has no explicit mechanism for comparing complete candidate repairs.

The correction metrics were saved to:

```
artifacts/correction/metrics.json  
artifacts/correction/model.pt
```

## 6. Training curves, convergence, and regularisation

The training curves for the binary detection model are saved in:

```
artifacts/detection/training_curves.png
```




Figure 2: Detection training curves

The detection model converged fairly cleanly. The first epoch was close to chance, with training accuracy 0.514 and development accuracy 0.510. By epoch 8, training accuracy had reached 0.754, while development accuracy reached 0.860. The development loss also decreased from 0.691 to 0.351.

This run does not show strong over-fitting. In fact, the development score is higher than the training score at the end of training. This can happen because the training metric is accumulated while the model is still being updated within the epoch, whereas the development metric is computed after the full epoch with dropout disabled. More importantly, the development loss is still decreasing at epoch 8, so the model has probably not reached its best possible in-distribution performance.

The correction curves are saved in:

```
artifacts/correction/training_curves.png
```

For correction, the loss also decreases steadily: training loss moves from 1.011 to 0.344, and development loss moves from 0.773 to 0.303. However, the task-level behaviour is less satisfactory. Development token accuracy starts very high, around 0.976, because most labels are OK, then decreases to around 0.913 as the weighted loss encourages the model to predict more edit labels. Exact-match accuracy remains almost flat, moving only from 0.490 to 0.516.

This means that the correction model is learning something about rare edit labels, but not enough to produce globally correct repairs. The loss curve alone is therefore not sufficient: for this task, exact-match accuracy is the more meaningful metric.

The regularisation strategies used in both models were:

- dropout 0.1 inside the Transformer encoder;
- AdamW weight decay 0.01;
- gradient clipping with maximum norm 1.0;
- small model size, with 2 layers and hidden dimension 64.

For correction, I also used inverse-square-root class weighting to compensate for the dominance of the OK label. A stronger inverse-frequency weighting was tested, but it over-corrected: the model began editing correct strings and exact-match accuracy collapsed. The final setting is therefore a compromise between learning rare edits and preserving already-correct sequences.

## 6. Out-of-distribution generalisation

### 7. Direct evaluation on deeper nesting

I implemented OOD evaluation in `lab4_ood.py`. The script loads the trained binary detection model from `artifacts/detection/model.pt` and evaluates it on strings with exact maximum nesting depth  $n = 5, 6, 7$ .

One implementation detail is important: the detection model was trained with `max_len = 80`, including `[CLS]` and `[SEP]`. To avoid truncating the raw bracket strings during OOD evaluation, I used raw lengths from 40 to 78 rather than 40 to 80.

Another data-generation detail matters for the OOD split. The generator from the handout only guarantees maximum depth  $\leq n$ . For experiments described as exact depth  $n$ , I explicitly computed `max_depth(string)` and rejected samples whose actual maximum depth was not exactly  $n$ . This is used for the OOD evaluations at depths 5, 6, and 7, and for the depth-5 fine-tuning experiment. In contrast, the length-augmentation experiment deliberately keeps the training bound at  $n \leq 4$ , so it may include depths 1-4 but not deeper strings.

The evaluation command was:

```
python3 lab4_ood.py \
  --examples-per-depth 1000 \
  --output-dir artifacts/ood
```

The overall OOD accuracy was:

| Split                        | Accuracy |
|------------------------------|----------|
| In-distribution test         | 0.852    |
| OOD test, $n$ in $\{5,6,7\}$ | 0.556    |

Accuracy as a function of maximum nesting depth:

| Depth | Accuracy |
|-------|----------|
| 5     | 0.560    |
| 6     | 0.570    |
| 7     | 0.539    |

Accuracy as a function of sequence length:

| Length bin | Accuracy |
|------------|----------|
| 40-49      | 0.629    |
| 50-59      | 0.639    |
| 60-69      | 0.516    |
| 70-78      | 0.464    |

The plots are saved in:

```
artifacts/ood/accuracy_by_depth.png
artifacts/ood/accuracy_by_length.png
artifacts/ood/metrics.json
```




Figure 3: OOD accuracy by depth




Figure 4: OOD accuracy by length

The model clearly does not generalise robustly out of distribution. Accuracy drops from 0.852 in distribution to 0.556 on deeper and longer OOD strings. The degradation is not neatly linear in depth: performance is similar for depths 5 and 6, then slightly worse for depth 7. The length trend is more revealing. Accuracy is still above chance for lengths 40-59, but falls near chance for 60-69 and below chance for 70-78.

I read this as a failure of length-invariant generalisation. The model probably relies on patterns that are useful within the training range, such as local bracket configurations, length cues, and shallow-depth regularities. Once sequences become substantially longer, those heuristics stop being reliable. The performance drop is therefore closer to an abrupt length-driven failure than to a smooth, rule-like degradation with nesting depth.

## 8. Pushdown automaton baseline

I implemented the deterministic pushdown automaton baseline in `lab4_pda_baseline.py`. The algorithm is the exact Dyck membership test:

1. Scan the sequence from left to right.
2. Push each opener onto a stack.
3. When a closer is read, reject if the stack is empty.
4. Otherwise pop the top opener and check that it matches the closer type.
5. Accept if and only if the stack is empty at the end.

The command was:

```
python3 lab4_pda_baseline.py \
  --id-size 1000 \
  --ood-size-per-depth 1000 \
  --output-dir artifacts/pda
```

The PDA results were:

| Split           | Accuracy |
|-----------------|----------|
| In-distribution | 1.000    |
| OOD depth 5     | 1.000    |
| OOD depth 6     | 1.000    |
| OOD depth 7     | 1.000    |
| OOD overall     | 1.000    |

The baseline also achieves perfect accuracy for every error type E1-E4. The metrics are saved in:

`artifacts/pda/metrics.json`

I also added a parity-only detection baseline. It predicts “error” if the raw bracket string has odd length and “no error” otherwise. This baseline is intentionally simple: it checks whether the Transformer is being compared against a cheap shortcut. In this data setup, E1, E2, and E4 change length parity, while E3 type mismatches do not.

The comparison with the Transformer and PDA is:

| Model                | In-distribution accuracy | OOD accuracy |
|----------------------|--------------------------|--------------|
| Parity-only baseline | 0.875                    | 0.875        |
| Transformer detector | 0.852                    | 0.556        |
| Pushdown automaton   | 1.000                    | 1.000        |

Parity-only accuracy by error type is:

| Error type        | Accuracy |
|-------------------|----------|
| No error          | 1.000    |
| E1 Missing closer | 1.000    |

| Error type         | Accuracy |
|--------------------|----------|
| E2 Spurious opener | 1.000    |
| E3 Type mismatch   | 0.000    |
| E4 Premature close | 1.000    |

This is a useful sanity check. The parity baseline is very strong because three of the four corruption types change string length parity. The Transformer does not even match this shortcut on the OOD split, so its poor OOD behaviour cannot be explained as a robust use of parity alone. It also explains why E3 is the important non-parity error type: detecting E3 requires more than odd/even length.

This reveals an important difference between solving the task extensionally and solving it algorithmically. The PDA has an explicit stack, so the same computation applies regardless of sequence length or nesting depth. Its representation is symbolic and invariant to the training distribution.

The Transformer, by contrast, learns distributed vector representations from bounded data. Even though self-attention can compare distant tokens, the model is not forced to implement a stack. Its poor OOD performance suggests that the learned representation captures useful in-distribution correlations, but not the exact recursive rule for Dyck membership. In other words, the Transformer approximates the language well inside the training regime, while the PDA recognises the language by construction.

## 9. Fine-tuning on 500 examples with depth 5

I implemented the fine-tuning experiment in `lab4_finetune_ood.py`. Starting from the trained binary detector, I fine-tuned on 500 examples with exact maximum nesting depth `n = 5`, then evaluated on fresh OOD examples with depths 5, 6, and 7.

The command was:

```
python3 lab4_finetune_ood.py \
  --finetune-size 500 \
  --eval-size-per-depth 1000 \
  --epochs 10 \
  --output-dir artifacts/finetune_ood
```

The results were:

The absolute “before” numbers in this table are not identical to the OOD numbers in question 7 because each script generated a fresh synthetic OOD evaluation sample, although with fixed random seeds. Comparisons should therefore be read within each table. A cleaner follow-up would persist one shared OOD JSONL split and reuse it across all scripts.

| Depth | Before fine-tuning | After fine-tuning | Delta  |
|-------|--------------------|-------------------|--------|
| 5     | 0.598              | 0.615             | +0.017 |
| 6     | 0.553              | 0.613             | +0.060 |
| 7     | 0.568              | 0.635             | +0.067 |

The metrics and plot are saved in:

```
artifacts/finetune_ood/metrics.json
artifacts/finetune_ood/before_after_accuracy.png
```

Fine-tuning on `n = 5` improves performance not only on depth 5 but also on depths 6 and 7. There is some transfer from the new examples, but it is modest: even after fine-tuning, accuracy remains around 0.61–0.64, far below the PDA baseline.

I would call this weak transfer rather than strong compositional generalisation. If the model had learned the true recursive rule from the `n = 5` examples, I would expect a larger and more stable gain on `n = 6` and `n = 7`. Instead, the model seems to adapt to some broader properties of the OOD regime, such as longer sequences and deeper local bracket patterns, without acquiring a fully depth-invariant stack algorithm.

## 7. Attention analysis

### 10. Average attention matrices and bracket matching heads

I implemented attention extraction in `lab4_attention.py`. Since PyTorch's standard `TransformerEncoderLayer` does not return attention weights by default, the script manually replays the encoder forward pass and calls each layer's `self_attn` module with `average_attn_weights=False`. This gives one attention matrix per layer and per head.

The analysis used 100 correct Dyck strings of length 40:

```
python3 lab4_attention.py \  
  --num-examples 100 \  
  --length 40 \  
  --output-dir artifacts/attention
```

The attention plots are saved as:

```
artifacts/attention/layer_0_head_0.png  
artifacts/attention/layer_0_head_1.png  
artifacts/attention/layer_0_head_2.png  
artifacts/attention/layer_0_head_3.png  
artifacts/attention/layer_1_head_0.png  
artifacts/attention/layer_1_head_1.png  
artifacts/attention/layer_1_head_2.png  
artifacts/attention/layer_1_head_3.png
```

For each matched pair  $(i, j)$ , I computed:

```
alpha_i->j = A[i, j]  
alpha_j->i = A[j, i]
```

where  $i$  is the opener position,  $j$  is the matched closer position, and  $A$  is the attention matrix for a given head. Because the analysed sequences have length 40 plus `[CLS]` and `[SEP]`, a roughly uniform attention baseline is approximately:

$$1 / 42 = 0.0238$$

The strongest heads by bidirectional matched-pair attention were:

| Layer | Head | Mean matched-pair attention |
|-------|------|-----------------------------|
| 0     | 3    | 0.0284                      |
| 0     | 1    | 0.0240                      |
| 1     | 0    | 0.0237                      |
| 0     | 2    | 0.0228                      |
| 1     | 3    | 0.0219                      |

Detailed opener-to-closer and closer-to-opener scores for the strongest head were:

| Layer | Head | alpha_i->j |                | alpha_j->i |                |
|-------|------|------------|----------------|------------|----------------|
|       |      | mean       | alpha_i->j std | mean       | alpha_j->i std |
| 0     | 3    | 0.0404     | 0.0143         | 0.0164     | 0.0144         |

The results do not show a clean bracket-matching head. Layer 0 head 3 attends from openers to their matching closers slightly more than uniform attention, but the effect is modest and asymmetric: closers do not strongly attend back to their openers. Most other heads are at or below the uniform baseline.

This lines up with the OOD results: the model can solve many in-distribution examples, but the attention plots do not expose an explicit bracket-pairing mechanism. The useful computation may be distributed across heads and layers, or the model may mainly rely on shallower statistical cues rather than systematic matching.




Figure 5: Attention layer 0 head 3

The attention metrics are saved in:

`artifacts/attention/metrics.json`

## 11. Attention patterns on erroneous strings

I implemented the erroneous-string attention analysis in `lab4_attention_errors.py`. The script generates corrupted strings, records the corrupted position, extracts attention for the same head that had the highest matched-pair score in the previous section, and compares the corrupted position to a random control position in the same sequence.

The command was:

```
python3 lab4_attention_errors.py \
  --examples-per-type 25 \
  --length 40 \
  --layer 0 \
  --head 3 \
  --output-dir artifacts/attention_errors
```

For E1, the missing closer is not an observed token, so I anchored the error position to the token immediately before the deletion point. For E2, E3, and E4, the error position is the inserted or replaced token.

The script saves local attention windows centred on the corrupted position:

```
artifacts/attention_errors/E1_missing_closer_local_window.png
artifacts/attention_errors/E2_spurious_opener_local_window.png
artifacts/attention_errors/E3_type_mismatch_local_window.png
artifacts/attention_errors/E4_premature_close_local_window.png
```

The numerical summary for layer 0 head 3 was:

| Error type         | Error incoming | Control incoming | [CLS] -> error | [CLS] -> control | Error local mass | Control local mass |
|--------------------|----------------|------------------|----------------|------------------|------------------|--------------------|
| E1 Missing closer  | 0.0237         | 0.0226           | 0.0224         | 0.0231           | 0.161            | 0.155              |
| E2 Spurious opener | 0.0226         | 0.0280           | 0.0166         | 0.0290           | 0.156            | 0.164              |
| E3 Type mismatch   | 0.0215         | 0.0223           | 0.0248         | 0.0215           | 0.137            | 0.160              |
| E4 Premature close | 0.0269         | 0.0213           | 0.0294         | 0.0231           | 0.159            | 0.152              |

The same weak matching head does not become a clean error-localisation head on erroneous strings. Most differences between corrupted and control positions are small. E4 is the clearest exception: the premature closer receives higher incoming attention and more attention from [CLS] than the control position. A cautious reading is that the model sometimes routes information from globally salient anomalies toward the classification token.

However, the effect is not strong enough to support a claim that attention directly implements error localisation. E2 even shows the opposite tendency: the corrupted opener receives less [CLS] attention than the control position. E3 has slightly higher [CLS] -> error attention, but lower local attention mass around the corrupted position.

My hypothesis is therefore cautious: attention may participate in localising some errors, especially premature closes that violate prefix well-formedness, but the model does not use a stable, interpretable attention pattern that corresponds to the symbolic stack algorithm. Error evidence is likely distributed across several heads and representations, rather than concentrated in one obvious matching head.

The metrics are saved in:

`artifacts/attention_errors/metrics.json`

## 12. Bonus: attention rollout to [CLS]

I implemented attention rollout in `lab4_attention_rollout.py`, following the usual idea of combining attention matrices across layers. For each layer, I averaged over heads, added an identity matrix to account for residual connections, row-normalised, and multiplied the resulting matrices across layers. I then inspected the rollout row for [CLS], which gives an approximate measure of how much information from each input token flows into the classification token.

The command was:

```
python3 lab4_attention_rollout.py \
  --examples-per-type 50 \
  --length 40 \
  --output-dir artifacts/attention_rollout
```

The results compare the rollout weight of the corrupted position against ordinary non-special tokens in the same sequence:

| Error type         | Error weight | Control mean | Error/control | Error percentile |
|--------------------|--------------|--------------|---------------|------------------|
| E1 Missing closer  | 0.0181       | 0.0189       | 0.963         | 0.485            |
| E2 Spurious opener | 0.0170       | 0.0179       | 0.952         | 0.480            |
| E3 Type mismatch   | 0.0216       | 0.0183       | 1.184         | 0.693            |
| E4 Premature close | 0.0201       | 0.0179       | 1.125         | 0.666            |

The rollout plots are saved in:

```
artifacts/attention_rollout/E1_missing_closer_rollout.png
artifacts/attention_rollout/E2_spurious_opener_rollout.png
artifacts/attention_rollout/E3_type_mismatch_rollout.png
artifacts/attention_rollout/E4_premature_close_rollout.png
artifacts/attention_rollout/metrics.json
```

The result is mixed. Tokens involved in E3 and E4 receive somewhat larger rollout weights than ordinary tokens: around 18% higher for type mismatches and 12% higher for premature closes. They also sit around the 67th-69th percentile of token rollout weights. In this run, visible closing-bracket anomalies therefore receive slightly more information flow into [CLS].

However, E1 and E2 do not show this effect. For E1, the error is an absence rather than an observed token, so anchoring the missing closer to a neighbouring token is necessarily imperfect. For E2, the spurious opener does not receive elevated rollout weight in this run. Thus, attention rollout gives weak evidence that some error tokens become more salient to [CLS], but the effect is neither universal nor strong enough to count as a reliable localisation mechanism.

## 8. Probing classifiers

### 13. Global depth probe

I implemented the global depth probe in `lab4_global_probe.py`. The script freezes the trained Transformer detector, extracts the final-layer [CLS] representation for correct Dyck strings, and trains two linear probes:

1. a linear regression probe predicting maximum nesting depth as a scalar;
2. a multinomial logistic-regression probe predicting depth as a discrete label in  $\{1, \dots, 7\}$ .

The command was:

```
python3 lab4_global_probe.py \  
  --examples-per-depth 500 \  
  --output-dir artifacts/global_probe
```

The probe dataset was balanced across depths 1-7, with 500 examples per depth. I used a 70/30 train/test split for the probes. The Transformer itself was frozen.

The results were:

| Probe                 | Metric   | Value |
|-----------------------|----------|-------|
| Linear regression     | $R^2$    | 0.415 |
| Linear classification | Accuracy | 0.270 |

Chance accuracy for the 7-way classifier is approximately:

$$1 / 7 = 0.143$$

Per-depth classification accuracy was:

| Depth | Accuracy |
|-------|----------|
| 1     | 0.549    |
| 2     | 0.234    |
| 3     | 0.231    |
| 4     | 0.118    |
| 5     | 0.147    |
| 6     | 0.196    |
| 7     | 0.385    |

The metrics are saved in:

`artifacts/global_probe/metrics.json`

The regression result shows that maximum nesting depth is partially linearly recoverable from the final [CLS] representation. An  $R^2$  of 0.415 is well above zero, so [CLS] does contain some global depth information.

However, the classification result is weak. Accuracy 0.270 is above chance, but far from a clean encoding of the seven depth classes. The classifier does especially poorly for middle depths 4-6. The representation contains a rough signal correlated with depth, but not a sharply separated linear code for exact maximum nesting depth.

This also fits the OOD results. The Transformer appears to encode some distributional features related to depth, but not a precise symbolic depth variable that would support robust generalisation to deeper nesting.

#### 14. Local depth probe

I implemented the local depth probe in `lab4_local_probe.py`. For each token position `t`, the target is the current nesting depth after reading the prefix up to that token:

```
depth_t = number of unmatched openers in s_1 ... s_t
```

The script freezes the Transformer, extracts token representations from the embedding layer and from each Transformer layer, and trains a linear ridge-regression probe for each layer.

The command was:

```
python3 lab4_local_probe.py \  
  --examples-per-depth 300 \  
  --output-dir artifacts/local_probe
```

This produced 100,632 token-level examples. The results were:

| Representation       | Probe $R^2$ |
|----------------------|-------------|
| Embedding + position | 0.139       |
| Layer 1              | 0.160       |
| Layer 2              | 0.160       |

The plot is saved in:

artifacts/local\_probe/r2\_by\_layer.png  
artifacts/local\_probe/metrics.json




Figure 6: Local depth probe R2 by layer

The best local depth probe is at layer 2, but the difference between layer 1 and layer 2 is tiny. The overall  $R^2$  is low. This means that current stack depth is only weakly linearly recoverable from individual token representations.

This result is informative because local depth is exactly the kind of quantity that a stack-based Dyck recogniser would maintain explicitly. If the Transformer had learned a clean stack-like algorithm, we might expect local depth to become increasingly linearly decodable in later layers. Instead, the probe improves only slightly over the positional/token embedding baseline. The model therefore does not appear to store local depth as a simple linear feature of each contextual token representation.

## 15. Probes on correct versus erroneous strings

I implemented the erroneous-string probing experiment in `lab4_probe_errors.py`. The probe is trained on correct strings only, using the final token representations from layer 2. It is then evaluated on:

1. held-out correct strings;
2. corrupted strings with known error positions.

For erroneous strings, I used the corrupted prefix balance as the target depth signal. This means that an inserted premature closer can make the balance drop too early, and a missing closer can make the later balance too high. This is not a perfect semantic target for “intended” repair, but it tests whether the model representation tracks the actually observed prefix structure.

The command was:

```
python3 lab4_probe_errors.py \
  --train-examples-per-depth 300 \
  --test-examples-per-depth 100 \
  --error-examples-per-type 100 \
  --layer-index 2 \
  --output-dir artifacts/probe_errors
```

The probe performance on held-out correct strings was:

| Metric              | Value |
|---------------------|-------|
| $R^2$               | 0.153 |
| Mean absolute error | 1.285 |

On erroneous strings, the global  $R^2$  against corrupted prefix balance was negative:

| Metric | Value  |
|--------|--------|
| $R^2$  | -0.218 |

Mean absolute error before, at, and after the corrupted position:

| Error type         | Before | At error | After |
|--------------------|--------|----------|-------|
| E1 Missing closer  | 1.073  | 0.949    | 0.756 |
| E2 Spurious opener | 1.126  | 0.876    | 0.722 |
| E3 Type mismatch   | 1.039  | 1.217    | 1.212 |
| E4 Premature close | 0.992  | 1.673    | 1.739 |

The E3 and E4 patterns are the most interpretable. For type mismatches and premature closes, the probe error increases at and after the corrupted token. E4 is especially clear: the mean absolute error rises from 0.992 before the error to 1.673 at the error and 1.739 after it. After the prefix structure is corrupted, the representation becomes less compatible with the depth probe.

For E1 and E2, the pattern is less useful. The error is either an absence or an extra opener, and the local depth probe does not show a clean breakdown after the corruption. That is not surprising given that the local depth probe was already weak on correct strings.

I also tried using the probe residual as an unsupervised error-localisation score. For each erroneous sequence, I ranked positions by absolute probe error and checked whether the true error position had the largest residual.

| Error type         | Mean rank of true error | Top-1 localisation accuracy |
|--------------------|-------------------------|-----------------------------|
| E1 Missing closer  | 28.93                   | 0.020                       |
| E2 Spurious opener | 31.93                   | 0.030                       |
| E3 Type mismatch   | 25.93                   | 0.050                       |
| E4 Premature close | 26.46                   | 0.010                       |

This localisation attempt fails. The true error position is rarely the position with the largest probe residual. The reason is straightforward: local depth is not strongly linearly encoded in the first place, so residuals from this weak probe are too noisy to serve as a reliable localisation signal.

The metrics are saved in:

artifacts/probe\_errors/metrics.json

## 16. Bonus: layer-wise probes and progressive computation

I implemented the layer-wise global probe in `lab4_layerwise_global_probe.py`. This extends the global depth analysis by probing [CLS] after the embedding layer, after layer 1, and after layer 2. The local depth layer-wise results come from `lab4_local_probe.py`.

The global probe command was:

```
python3 lab4_layerwise_global_probe.py \  
--examples-per-depth 500 \  
--output-dir artifacts/layerwise_global_probe
```

Global maximum-depth probe results:

| Representation | Regression $R^2$ | Classification accuracy |
|----------------|------------------|-------------------------|
| Embedding      | -0.001           | 0.132                   |
| Layer 1        | 0.391            | 0.255                   |
| Layer 2        | 0.426            | 0.295                   |

Local current-depth probe results:

| Representation | Local depth $R^2$ |
|----------------|-------------------|
| Embedding      | 0.139             |
| Layer 1        | 0.160             |
| Layer 2        | 0.160             |

The plots are saved in:

```
artifacts/layerwise_global_probe/global_r2_by_layer.png  
artifacts/layerwise_global_probe/global_accuracy_by_layer.png  
artifacts/local_probe/r2_by_layer.png
```

The global depth probes show a progressive pattern. The embedding-level [CLS] representation contains essentially no useful global depth information: regression  $R^2$  is around zero and classification accuracy is close to chance. After the first Transformer layer, depth becomes much more linearly decodable. After the second layer, it improves slightly further.

The local depth probes show a much weaker progression. Token-level current depth is already weakly predictable from token and position embeddings, and the Transformer layers improve it only slightly. I do not see evidence here for an explicit linear stack-depth variable at every token.

Overall, the results support a limited version of progressive computation. The model seems to accumulate some global information into [CLS] across layers, but this is not matched by a clean layer-by-layer construction of local stack state. In other words, the network develops a rough global summary useful for classification, rather than a transparent symbolic computation of Dyck depth at each position.

## 9. Open discussion

### 17. Rule learning versus statistical pattern memorisation

Based on the experiments above, I would not conclude that this Transformer has learned the full Dyck membership rule. It has learned useful in-distribution regularities, but the evidence points more toward bounded statistical pattern learning than toward a robust stack-like algorithm.

The strongest evidence comes from the OOD evaluation. In distribution, the binary detector achieves accuracy 0.852. On deeper and longer OOD strings, accuracy drops to 0.556. The performance drop is especially severe for long strings: the model reaches only 0.464 accuracy on lengths 70-78. A true Dyck recogniser should not degrade in this way, because the membership rule is the same at all depths and lengths.




Figure 7: Global depth probe R2 by layer




Figure 8: Global depth probe accuracy by layer




Figure 9: Local depth probe R2 by layer

The PDA comparison makes the contrast explicit:

| Model                | In-distribution accuracy | OOD accuracy |
|----------------------|--------------------------|--------------|
| Transformer detector | 0.852                    | 0.556        |
| Pushdown automaton   | 1.000                    | 1.000        |

The PDA uses a stack, so its computation is depth-invariant. The Transformer does not show this kind of invariance. Fine-tuning on 500 examples with depth 5 improves depths 6 and 7 slightly, but final accuracies remain around 0.61–0.64. I read that as weak transfer to the broader OOD regime, not acquisition of the exact recursive rule.

The attention analysis points in the same direction. If the model had learned a transparent bracket-matching algorithm, we might expect one or more attention heads to strongly connect matched openers and closers. Instead, the strongest matched-pair head has only modest attention above the uniform baseline, and the effect is asymmetric. Attention rollout gives some extra salience to E3 and E4 error tokens, but not to E1 and E2. Thus, attention does not reveal a stable symbolic matching mechanism.

The probing results also support a cautious conclusion. Maximum depth is partially recoverable from [CLS], with global regression  $R^2$  around 0.415, and the layer-wise global probe shows that this information becomes more available after the Transformer layers. However, local depth is only weakly recoverable from token representations, with  $R^2$  around 0.160. Since a stack-based recogniser would naturally maintain local depth and opener type information, the weak local probe result is evidence against a clean internal stack representation.

The correction experiment is another warning sign. Token-level correction accuracy is high, but exact-match accuracy is only 0.509, barely above the “leave correct strings alone” baseline. The model does not reliably construct globally valid repairs. This is what we would expect if it has learned local edit tendencies rather than a global grammar-constrained repair process.

My conclusion is therefore:

The Transformer has learned useful bounded-depth heuristics for Dyck-like strings, but not the exact Dyck membership rule.

These heuristics include local bracket configurations, length-related cues, rough global depth signals in [CLS], and some sensitivity to visible closing-bracket anomalies. They are enough for reasonable in-distribution performance, but they do not compose reliably when depth and length move outside the training regime.

## 18. Length-based data augmentation for OOD generalisation

I implemented a length-based data augmentation experiment. The motivation comes from the OOD results in question 7: performance degraded most clearly as sequence length increased. The original detector was trained on lengths 4-40, while the OOD split uses lengths around 40-78 in my non-truncated evaluation. Therefore, I tried exposing the model to longer strings while keeping the same training-depth bound  $n \leq 4$ .

I first tried training from scratch with length augmentation in `lab4_length_aug.py`, using lengths 4-78. This did not work well under the same small training budget: the model stayed near chance. Longer strings make the task harder, so starting from scratch with the same 10k examples and 8 epochs was not enough.

The successful variant was length-augmentation fine-tuning, implemented in `lab4_length_aug_finetune.py`. I started from the already trained detector and fine-tuned it on 10,000 examples with:

```
depth:  n <= 4
length: 4-78
```

The command was:

```
python3 lab4_length_aug_finetune.py \
  --finetune-size 10000 \
  --dev-size 1000 \
  --ood-size-per-depth 1000 \
  --epochs 8 \
  --output-dir artifacts/length_aug_finetune
```

The OOD results before and after length-augmentation fine-tuning were:

As in question 9, these “before” values come from a fresh synthetic OOD evaluation sample generated inside this script. They are close to, but not numerically identical with, the question 7 OOD table. The before/after comparison below should be interpreted within this table.

| Depth   | Before | After | Delta  |
|---------|--------|-------|--------|
| 5       | 0.584  | 0.756 | +0.172 |
| 6       | 0.560  | 0.765 | +0.205 |
| 7       | 0.567  | 0.780 | +0.213 |
| Overall | 0.570  | 0.767 | +0.197 |

The fine-tuning development accuracy on length-augmented data reached 0.844 by epoch 8. The metrics and curve are saved in:

```
artifacts/length_aug_finetune/metrics.json
artifacts/length_aug_finetune/training_curves.png
```

This modification substantially improves OOD performance, especially for depths 6 and 7. That matches the earlier observation that much of the OOD failure was length-driven. Once the model sees longer sequences during fine-tuning, it becomes much less brittle on the long OOD examples.

I still would not treat this as proof of rule learning. The model is not perfect, and the PDA remains at 1.000 accuracy. Length augmentation teaches the model to handle a broader input distribution, but it does not show that the model has learned an unbounded stack algorithm. The more conservative reading is improved interpolation over a wider range of lengths, with some useful transfer to deeper nesting.




Figure 10: Length augmentation fine-tuning curves

## 19. Relevance to syntactic capabilities of large language models

The Dyck task is informative about the syntactic capabilities of large language models, but only in a limited and carefully qualified way.

It is informative because Dyck languages isolate one important component of syntax: hierarchical nesting. Natural language contains many dependencies that are not purely local, such as embedded clauses, agreement across intervening material, and recursively nested constituents. A model that cannot handle even a clean synthetic nesting task should make us cautious about claims that it has learned fully systematic syntax. In this sense, Dyck languages are useful stress tests for structural generalisation.

The experiments here show a pattern that is relevant to LLMs more broadly: high in-distribution performance does not necessarily imply rule-like generalisation. The Transformer detector works reasonably well on the training-like test set, but its OOD performance drops sharply when length and depth increase. Attention and probing analyses also fail to reveal a clean symbolic stack. The practical lesson is that neural models may solve syntactic-looking tasks using a mixture of shallow heuristics, distributional regularities, and partial structural representations.

However, there are important disanalogies between Dyck languages and natural language parsing.

First, Dyck syntax is deterministic, fully explicit, and noise-free. Every bracket has an unambiguous type, and well-formedness can be decided exactly by a stack. Natural language syntax is not like this. It contains ambiguity, optionality, ellipsis, long-distance dependencies, semantic constraints, discourse effects, and constructions whose analysis may depend on world knowledge. A sentence is not simply accepted or rejected by matching visible symbols.

Second, Dyck languages have no lexical semantics. In natural language, syntactic expectations are deeply entangled with word meaning, selectional preferences, morphology, and pragmatics. For example, verb argument structure and agreement both depend on lexical and grammatical features, not just on abstract brackets. A model may succeed or fail on natural language for reasons that are invisible in the Dyck setting.

Third, the training regimes are very different. The small Transformer in this lab is trained from scratch on tens of thousands of synthetic examples. Large language models are trained on massive corpora containing many overlapping cues: punctuation, formatting, code, natural language, parallel paraphrases, and repeated constructions. Their syntactic behaviour may reflect scale and data diversity in ways that this lab does not capture.

Fourth, natural language rarely requires unbounded centre embedding in practice. Humans also struggle with deeply nested sentences, and many natural-language dependencies are shallow or right-branching. Therefore, failure on very deep Dyck strings does not directly imply failure on ordinary sentence syntax. Conversely, success on Dyck strings would not by itself imply mastery of natural language grammar.

So the right conclusion is moderate. Dyck experiments are valuable because they provide a controlled test of hierarchical generalisation and allow precise interpretability analyses. They can reveal whether a Transformer learns

stack-like computations under idealised conditions. But they are not a complete model of natural language syntax. The results here suggest caution: a Transformer can look syntactically competent inside a familiar distribution while still lacking a robust, interpretable rule for the underlying formal language.